

CMSC 330: Organization of Programming Languages

Threads


Multiprocessing

- Description
 - Multiple processing units (or multiple cores)
 - From single microprocessor to large computer clusters
 - Can perform multiple tasks in parallel simultaneously

CMSC 330


2

Computation Abstractions



A computer

Processes vs. Threads




Processes do not share data

Threads share data within a process

So, What Is a Thread?

- **Conceptually:** it is a parallel computation occurring within a process
- **Implementation view:** it's a program counter and a stack. The heap and static area are shared among all threads
- All programs have at least one thread (main)

Implementation View



- Per-thread stack pointer (sp) and instruction pointer (ip)
 - Saved in memory when thread suspended
 - Copied to hardware sp/ip when thread resumes

Tradeoffs

- Threads can increase performance
 - Parallelism on multiprocessors
 - Concurrency of computation and I/O
- Natural fit for some programming patterns
 - Event processing
 - Simulations
- But increased complexity
 - Need to worry about safety, liveness, composition
- And higher resource usage

Programming Threads

- Thread creation is inexpensive
- Threads reside on same physical processor
- Threads share memory, resources
 - Except for local thread variables
- Shared-memory programming paradigm
 - Threads communicate via shared data
 - Synchronization used to avoid race conditions
- Limited scalability (tens of threads)

Programming Processes

- Process creation is expensive
 - Request to operating system
- Processes may reside on separate processors
- Processes do not share memory
- Message-passing programming paradigm
 - Messages using I/O streams, sockets, network, files
- Processes must cooperate to communicate
 - Actions performed to send and receive data
- Highly scalable (thousands of processors)


Threads in Programming Languages

- Threads are available in many languages
 - C, C++, Java, Ruby, OCaml,...
- In older languages (e.g., C and C++), threads are a platform-specific add-on
 - Not part of the language specification
 - Implemented as code libraries (e.g., pthreads)
- In newer languages (e.g., Java, Ruby), threads are part of the language specification
 - Not dependent on operating system
 - Can utilize special keywords, syntax


Concurrency

- A *concurrent* program is one that has multiple threads that may be active at the same time
 - Might run on one CPU
 - The CPU alternates between running different threads
 - The *scheduler* takes care of the details
 - Switching between threads might happen *at any time*
 - Might run *in parallel* on a *multiprocessor* machine
 - One with more than one CPU
 - May have multiple threads per CPU

Scheduling Example (1)




Scheduling Example (2)



CMSC 330

13

Thread Creation



CMSC 330

14

Thread Creation in Java

- To explicitly create a thread:
 - Instantiate a `Thread` object (an object of class `Thread` or a subclass of `Thread`)
 - Override its `run()` method
 - Invoke the object's `start()` method
 - This will start executing the `Thread`'s `run()` method concurrently with the current thread
 - A thread terminates when its `run()` method returns

CMSC 330

15

Example: Alarms

- Goal: let's set alarms that will be triggered in the future
 - Input: time `t` (seconds) and message `m`
 - Result: we'll see `m` printed after `t` seconds

CMSC 330

16

Example: Synchronous alarms

```
while (there is more input to read) {
    System.out.print("Alarm> ");

    // read user input into timeout and msg

    // wait (in seconds)
    try {
        Thread.sleep(timeout * 1000);
    } catch (InterruptedException e) { }
    System.out.println("(" + timeout + ") " + msg);
}
```

Making It Threaded (1)

```
public class AlarmThread extends Thread {

    private String msg = null;
    private int timeout = 0;

    public AlarmThread(String msg, int time) {
        this.msg = msg;
        this.timeout = time;
    }

    public void run() {
        try {
            Thread.sleep(timeout * 1000);
        } catch (InterruptedException e) { }
        System.out.println("(" + timeout + ") " + msg);
    }

}
```

Making It Threaded (2)

```
while (there is more input to read) {
    System.out.print("Alarm> ");

    // read user input into m and tm

    if (m != null) {
        // start alarm thread
        Thread t = new AlarmThread(m, tm);
        t.start();
    }
}
```

Alternative: The Runnable Interface

- Extending [Thread](#) prohibits a different superclass
- Instead implement [Runnable](#)
 - Declares that the class has a [void run\(\)](#) method
- Construct a [Thread](#) from the [Runnable](#)
 - Constructor [Thread\(Runnable target\)](#)
 - Constructor [Thread\(Runnable target, String name\)](#)
- This approach is preferred

Thread Example Revisited

```
public class AlarmRunnable implements Runnable {

    private String msg = null;
    private int timeout = 0;

    public AlarmRunnable(String msg, int time) {
        this.msg = msg;
        this.timeout = time;
    }

    public void run() {
        try {
            Thread.sleep(timeout * 1000);
        } catch (InterruptedException e) { }
        System.out.println("(" + timeout + ") " + msg);
    }
}
```

CMSC 330

21

Thread Example Revisited (2)

```
while (there is more input to read) {
    System.out.print("Alarm> ");

    // read user input into m and tm

    if (m != null) {
        // start alarm thread
        Thread t = new Thread(new
                                AlarmRunnable(m, tm));
        t.start();
    }
}
```

CMSC 330

22

Notes: Passing Parameters

- `run()` doesn't take parameters
- We "pass parameters" to the new thread by storing them as private fields
 - In the extended class
 - Or the `Runnable` object
 - Example: the time to wait and the message to print in the `AlarmThread` class

CMSC 330

23

Concurrency and Shared Data

- Concurrency is easy if threads don't interact
 - Each thread does its own thing, ignoring other threads
 - Typically, however, threads need to communicate with each other
- Communication is done by *sharing* data
 - In Java, different threads may access the heap simultaneously
 - But the scheduler might interleave threads arbitrarily
 - Problems can occur if we're not careful.

CMSC 330

24

Race Condition Example

```
public class Example extends Thread {  
  
    private static int count = 0; // shared state  
  
    public void run() {  
        int y = count;  
        count = y + 1;  
    }  
  
    public static void main(String args[]) {  
        Thread t1 = new Example();  
        Thread t2 = new Example();  
        t1.start();  
        t2.start();  
    }  
}
```

CMSC 330

25

Race Condition Example

```
static int count = 0;
```

```
t1.run() {  
    int y = count;  
    count = y + 1;  
}
```

Shared state count = 0

```
t2.run() {  
    int y = count;  
    count = y + 1;  
}
```

Start: both threads ready to run. Each will increment the global count.

CMSC 330

26

Race Condition Example

```
static int count = 0;
```

```
t1.run() {  
    int y = count;  
    count = y + 1;  
}
```

Shared state count = 0

y = 0



```
t2.run() {  
    int y = count;  
    count = y + 1;  
}
```

T1 executes, grabbing the global counter value into its own y.

CMSC 330

27

Race Condition Example

```
static int count = 0;
```

```
t1.run() {  
    int y = count;  
    count = y + 1;  
}
```

Shared state **count = 1**

y = 0



```
t2.run() {  
    int y = count;  
    count = y + 1;  
}
```

T1 executes again, storing its value of y + 1 into the counter.

CMSC 330

28

Race Condition Example

```
static int count = 0;
```

```
t1.run() {  
    int y = count;  
    count = y + 1;  
}
```

Shared state count = 1

y = 0

■ ■ ■

```
t2.run() {  
    int y = count;  
    count = y + 1;  
}
```

y = 1

T1 finishes. T2 executes, grabbing the global counter value into its own y.

Race Condition Example

```
static int count = 0;
```

```
t1.run() {  
    int y = count;  
    count = y + 1;  
}
```

Shared state count = 2

y = 0

■ ■ ■

```
t2.run() {  
    int y = count;  
    count = y + 1;  
}
```

y = 1

T2 executes, storing its incremented count value into the global counter.

But When it's Run Again?

- **count** could be 1 (make sure you see why)

Why?

- Different schedules can lead to different outcomes
 - This is a *race condition* or *data race*
- A thread was preempted in the middle of an operation
 - Reading and writing **count** was supposed to be *atomic*-to happen with no interference from other threads
 - But the schedule (interleaving of threads) that was chosen allowed atomicity to be violated
 - These bugs can be extremely hard to reproduce, and so hard to debug
 - Depends on what scheduler chose to do, which is hard to predict

Question

- If instead of

```
int y = count;
count = y + 1;
```
- We had written

```
- count++;
```
- Would the result be any different?
- Answer: NO!
 - Don't depend on your intuition about atomicity

Question

- If you run a program with a race condition, will you always get an unexpected result?
 - No! It depends on the scheduler, and on the other threads/processes/etc, that are running on the same CPU
- Race conditions are hard to find

Synchronization

- Refers to mechanisms allowing a programmer to control the execution order of some operations across different threads in a concurrent program.
- Different languages have adopted different mechanisms to allow the programmer to synchronize threads.
- Java has several mechanisms

Aspects of Synchronization

- Atomicity
 - Locking to obtain mutual exclusion
 - What we most often think about
- Visibility
 - Ensuring that changes to object fields made in one thread are seen in other threads
- Ordering
 - Ensuring that you aren't surprised by the order in which statements are executed

Java's "synchronized" keyword

- Every object has an implicit associated lock
- `synchronized` allows you to acquire an object's lock
 - `synchronized (obj) { body }`
 - Obtains the lock associated with `obj`
 - Executes `body`
 - Releases the lock when scope is exited, even in cases of exception or method return


Example

```
static Object o = new Object();

void f() throws Exception {
    synchronized (o) {
        FileInputStream f = new FileInputStream("file.txt");
        // do something with f
        f.close();
    }
}
```

- Lock associated with `o` acquired before body executed, released even if exception is thrown

Discussion



- An object and its associated lock are different!
 - Holding the lock on an object does not affect what you can do with that object in any way
 - Ex:

```
synchronized (o) {           // acquires lock named o
    o.f();                     // you can call o's methods
    o.x = 3;                   // you can read and write o's fields
}
```

Example: Synchronizing on this

```
class C {
    int count;

    void inc() {
        synchronized (this) {
            count++;
        }
    }
}
```

```
C c = new C();
```

```
Thread 1
c.inc();
```

```
Thread 2
c.inc();
```

- Does this program have a race condition?
 - No, both threads acquire lock on the same object before they access shared data

Example: Synchronizing on this (cont'd)

```
class C {
    int count;

    void inc() {
        synchronized (this) {
            count++;
        }
    }

    void dec() {
        synchronized (this) {
            count--;
        }
    }
}
```

```
C c = new C();
```

```
Thread 1
c.inc();
```

```
Thread 2
c.dec();
```

- Race condition?
 - No, threads acquire lock on the same object before they access shared data

CMSC 330

41

Example: Synchronizing on this (cont'd)

```
class C {
    int count;

    void inc() {
        synchronized (this) {
            count++;
        }
    }
}
```

```
C c1 = new C();
C c2 = new C();
```

```
Thread 1
c1.inc();
```

```
Thread 2
c2.inc();
```

- Does this program have a race condition?
 - No, threads acquire different locks, but they write to different objects, so that's ok

CMSC 330

42

Synchronized Methods

- Marking a method as synchronized is the same as synchronizing on this in body of the method
 - The following two programs are the same

```
class C {
    int count;

    void inc() {
        synchronized (this) {
            count++;
        }
    }
}
```

```
class C {
    int count;

    synchronized void inc() {
        count++;
    }
}
```

CMSC 330

43

Synchronized Methods (cont'd)

```
class C {
    int count;

    void inc() {
        synchronized (this) {
            count++;
        }
    }

    synchronized void dec() {
        count--;
    }
}
```

```
C c = new C();
```

```
Thread 1
c.inc();
```

```
Thread 2
c.dec();
```

- Race condition?
 - No, both acquire same lock

CMSC 330

44

Locks (Java 1.5)

```
interface Lock {
    void lock();
    void unlock();
    ... /* some more stuff also */
}

class ReentrantLock implements Lock { ... }
```

- Explicit **Lock** objects are the same as the implicit lock used by synchronized keyword
- Only one thread can hold a lock at once
 - Other threads that try to acquire it *block* (or become suspended) until the lock becomes available
 - **unlock()** allows lock to be acquired by different thread

CMSC 330

45

Avoiding Interference: Synchronization

```
public class Example extends Thread {

    private static int count = 0;
    static Lock lock = new ReentrantLock();

    public void run() {
        lock.lock();
        int y = count;
        count = y + 1;
        lock.unlock();
    }

    ...
}
```

***Lock**, for protecting the shared state*

***Acquires** the lock; Only succeeds if not held by another thread*

***Releases** the lock*

CMSC 330

46

Applying Synchronization

```
int count = 0;

t1.run() {
    lock.lock();
    int y = count;
    count = y + 1;
    lock.unlock();
}

t2.run() {
    lock.lock();
    int y = count;
    count = y + 1;
    lock.unlock();
}
```

- Trace this and be sure you understand why it fixes the race condition

CMSC 330

47

Synchronization Example (Java 1.4)

```
public class Example extends Thread {

    private static int count = 0;
    static Object value = new Object();

    public void run() {
        synchronized (value) {
            int y = count;
            count = y + 1;
        }
    }

    ...
}
```

***Value**, any Java object*

***Acquires** the lock associated w/ value; only succeeds if not held by another thread, otherwise blocks*

***Releases** the lock*

CMSC 330

48

Different Locks Don't Interact

```
static int count = 0;
static Lock l =
    new ReentrantLock();
static Lock m =
    new ReentrantLock();
```

```
void inc() {
    l.lock();
    count++;
    l.unlock();
}
```

```
void inc() {
    m.lock();
    count++;
    m.unlock();
}
```

- This program has a race condition
 - Threads only block if they try to acquire a lock held by another thread