

CMSC 330: Organization of Programming Languages

Threads, con't.

Parallelizable Applications of Interest

- Knowledge discovery: mine and analyze massive amounts of distributed data
 - Discovering social networks
 - Real-time, highly-accurate common operating picture, on small, power-constrained devices
- Simulations (games?)
- Data processing
 - NLP, vision, rendering, in real-time
- Commodity applications
 - Parallel testing, compilation, typesetting, ...

CMSC 330

2

Multithreading (Java threads, pthreads)

- + Portable, high degree of control
- Low-level and unstructured
 - Thread management, synchronization via locks and signals essentially manual
 - Blocking synchronization is not compositional, which inhibits nested parallelism
 - Easy to get wrong, hard to debug
 - Data races, deadlocks all too common

CMSC 330

3

Parallel Language Extensions

- MPI – expressive, portable, but
 - Hard to partition data and get good performance
 - Temptation is to hardcode data locations, number of processors
 - Hard to write the program correctly
 - Little relation to the sequential algorithm
- OpenMP, HPF – parallelizes certain code patterns (e.g., loops), but
 - Limited to built-in types (e.g., arrays)
 - Code patterns, scheduling policies brittle

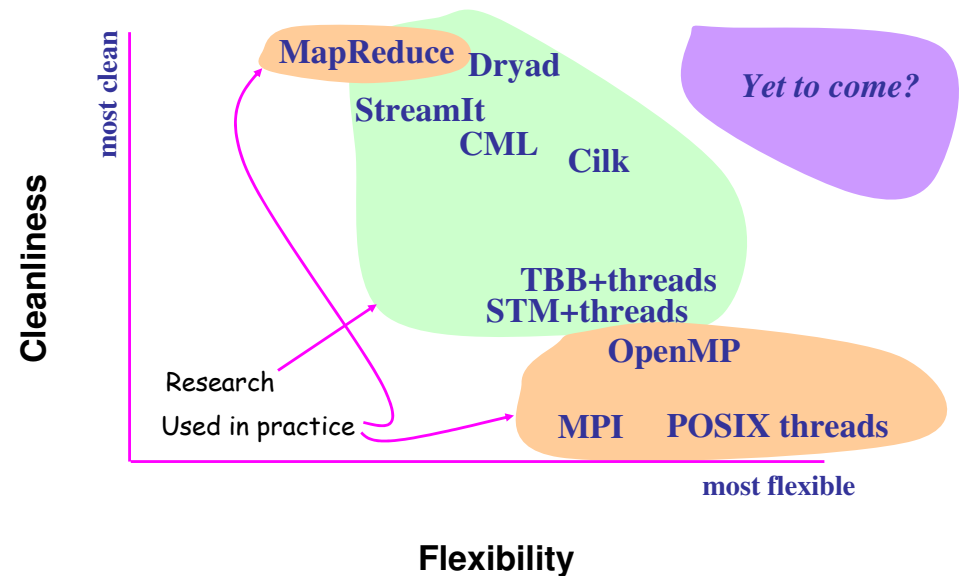
CMSC 330

4

Two Directions To A Solution

- Start with clean, but limited, languages/abstractions and generalize
 - MapReduce (Google)
 - StreamIt (MIT)
 - Cilk (MIT)
- Start with full-featured languages and add cleanliness
 - Software transactional memory
 - Static analyzers (Locksmith, Chord, ...)
 - Threaded Building Blocks (Intel)

Space of Solutions



Kinds of Parallelism

- Data parallelism
 - Can divide parts of the data between different tasks and perform the same action on each part in parallel
- Task parallelism
 - Different tasks running on the same data
- Hybrid data/task parallelism
 - A parallel pipeline of tasks, each of which might be data parallel
- Unstructured
 - Ad hoc combination of threads with no obvious top-level structure

MapReduce: Programming the Pipeline

- Pattern inspired by Lisp, ML, etc.
 - Many problems can be phrased this way
- Results in clean code
 - Easy to program / debug / maintain
 - Simple programming model
 - Nice retry/failure semantics
 - Efficient and portable
 - Easy to distribute across nodes

Map & Reduce in Lisp / Scheme

- (map *f list*)
Unary operator
- (map square '(1 2 3 4))
(1 4 9 16)
Binary operator
- (reduce + '(1 4 9 16) 0)
(+ 1 (+ 4 (+ 9 (+ 16 0))))
30
- (reduce + (map square '(1 2 3 4)) 0)

MapReduce a la Google

- map(key, val) is run on each item in set, emits new-key / new-val pairs
- reduce(key, vals) is run for each unique key emitted by map(), emits final output

Count Words in Documents

- Input consists of (url, contents) pairs
- map(key=url, val=contents):
 - For each word *w* in contents, emit (*w*, "1")
- reduce(key=word, values=uniq_counts):
 - Sum all '1's in values list
 - Emit result "(word, sum)"

Count, Illustrated


map(key=url, val=contents):

For each word *w* in contents, emit (*w*, "1")

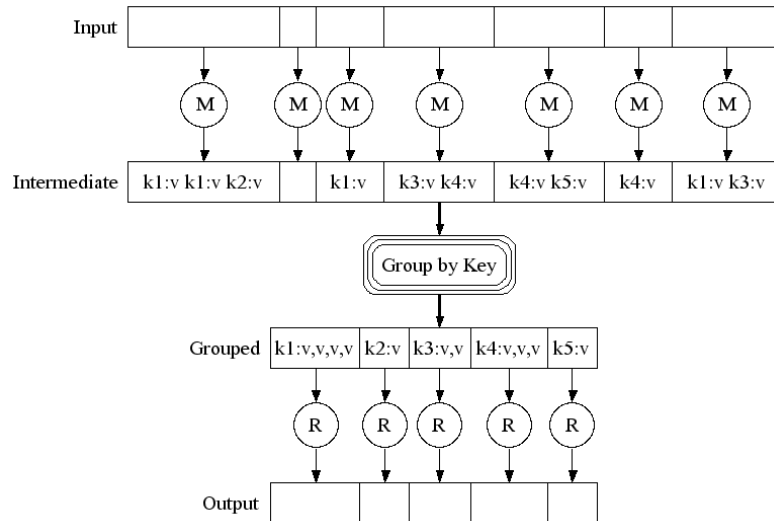
reduce(key=word, values=uniq_counts):

Sum all "1"s in values list

Emit result "(word, sum)"



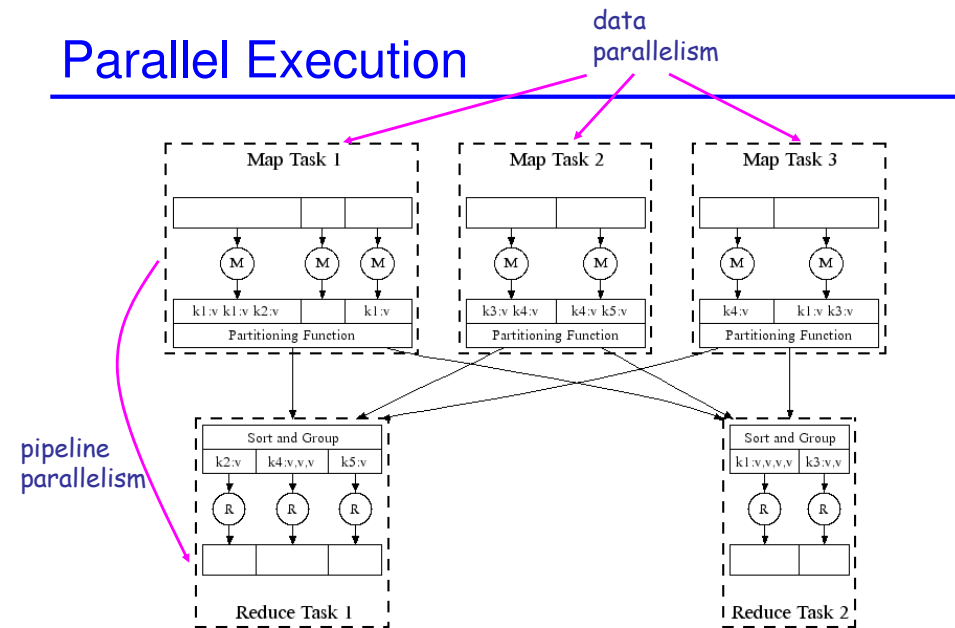
Execution



CMSC 330

13

Parallel Execution




Key: no implicit dependencies between map or reduce tasks

CMSC 330

14

Model is Widely Applicable

MapReduce Programs In Google Source Tree 2004



Example uses:

distributed grep
term-vector / host
document clustering
...

distributed sort
web access log stats
machine learning
...

web link-graph reversal
inverted index construction
statistical machine translation
...

CMSC 330

15

The Programming Model Is Key

- Simple control makes dependencies evident
 - Can automate scheduling of tasks and optimization
 - Map, reduce for different keys, embarrassingly parallel
 - Pipeline between mappers, reducers evident
- **map** and **reduce** are pure functions
 - Can rerun them to get the same answer
 - In the case of failure, or
 - To use idle resources toward faster completion
 - No worry about race conditions, deadlocks, etc. since there is no shared state

CMSC 330

16

Compare to Dedicated Supercomputers

- According to Wikipedia, in 2006 Google uses
 - 450,000 servers from 533 MHz Celeron to dual 1.4GHz Pentium III
 - 80GB drive per server, at least
 - 2-4GB memory per machine
 - Jobs processing 100 terabytes of distributed data
- More computing power than even the most powerful supercomputer