

## I Warmup

1. [4 points]: Which of the following is true about Address Resolution Protocol (ARP) and learning bridges?

(Circle ALL that apply)

- A. A learning bridge maintains state that maps IP addresses to hardware (MAC) addresses.
- B. A learning bridge maintains state that maps IP addresses to MAC addresses.
- C. A host's ARP table maintains state that maps IP addresses to hardware (MAC) addresses.
- D. A host's ARP table maintains state that maps hardware addresses to IP addresses.

2. [4 points]: Which of the following is true about DNS?

(Circle ALL that apply)

- A. A query for an A record may return multiple IP addresses in the response.
- B. A query for an NS record may return multiple IP addresses in the response.
- C. A query for a MX record may return multiple IP addresses in the response.
- D. A short TTL on an A record reply may run the risk of increasing traffic at the root nameserver.
- E. None of the above.

3. [2 points]: Which of the following most accurately describes the *most common* uses for eBGP, iBGP, and IGP?

(Circle the BEST answer)

- A. eBGP is used between ASes for external destinations, iBGP is used within an AS for external destinations, and IGP is used within an AS for destinations within an AS.
- B. eBGP is used within an AS for external destinations, iBGP is used between ASes for external destinations, and IGP is used within an AS for internal destinations.
- C. eBGP is used between ASes for external destinations, iBGP is used within an AS for internal destinations, and IGP is used within an AS for external destinations.
- D. None of the above

Name:

## II Potpourri

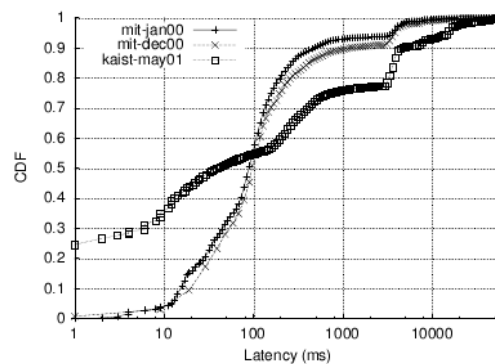
4. [4 points]: Explain the difference between a link-state routing protocol and a distance vector protocol.

(Answer legibly in the space below.)

5. [4 points]: Consider the graph below, which shows a cumulative distribution function (CDF) of the DNS response times for DNS lookup latencies for three data traces: (1) two different traces at a resolver at MIT in Cambridge, MA; (2) one trace at a resolver in Korea (*kaist*).<sup>1</sup>

- A. Which resolver has more queries that take more than a second to resolve?
- B. Why might the two traces have different latency distributions? (In other words, why might one resolver take longer to resolve queries than another?) There are several possible reasons; give one. (*Hint*: Think about geography.)

(Answer legibly in the space below.)



<sup>1</sup>In case you don't know, the way to read a CDF is as follows: a point  $(x, y)$  means that for that distribution, the fraction  $y$  of the points in that distribution have value  $x$  or smaller. For example, slightly more than 30% of the queries in the *kaist* trace had a lookup latency of 10ms or less.

Name: